

Projectivisation

Marco Kuhlmann

Goal Implement (a variant of) the Eisner algorithm for projective dependency parsing and apply it to a practical problem in dependency parsing.

Preparations Watch the video lecture on the Eisner algorithm on the course website.

Report Submit the requested code, as well as a short reflection piece (ca. 150 words) about your experience.

Deadline 2024-02-16

Introduction

Many algorithms for syntactic dependency parsing, including the arc-standard algorithm, are restricted to projective dependency trees. At the same time, several dependency treebanks contain non-projective dependency trees. One solution to this problem is to transform the (possibly non-projective) trees of the original treebank into projective trees prior to training – a transformation known as *projectivisation*. Your task in this practical is to implement a projectiviser based on the Eisner algorithm.

Background

When projectivising a dependency tree, it is desirable to maintain as much of the structure of the original tree as possible. More specifically, linguistic considerations tell us that it makes sense to only allow changes to the source tree that move the head of a dependency arc $h \rightarrow d$ to one of the *ancestors* a of h . This is known as *lifting*. Applying lifting to a dependency tree makes the tree “more” projective, in the sense that the number of crossing arcs in the tree at most decreases (compared to the source tree). In the extreme case, all arcs of the source tree can be lifted to the root vertex,

in which case the target tree is guaranteed to be projective. More generally though, we are interested in computing a projective dependency tree that requires a *minimal* amount of lifting. One can prove that, for every dependency tree, there is a unique projective dependency tree that satisfies the minimal lifting property.

To compute a projective dependency tree that requires a minimal amount of lifting, we can use a variant of the Eisner algorithm. Given an arc $h \rightarrow d$ from a head vertex h to a dependent vertex d in the source tree, we define the *cost* of a potential arc $h' \rightarrow d$ in the target tree as the number of lifting operations it would take to re-attach d to h' . For example, if h' is the parent of the original head vertex, this cost is 1; if h' is the grandparent, the cost is 2, and so on. For arcs that would simply be copied from the source tree to the target tree ($h' = h$), the cost is zero. More generally, the cost is equal to the length of the (unique) path in the source tree from h' to h , if such a path exists. For vertices h' that are not ancestors of h in the source tree, we set the cost to positive infinity. We can then adapt the Eisner algorithm to find a projective dependency tree with *minimal cost*, rather than maximal score.

Problem specification

You will have to write a Python script that

- reads a treebank file in the CoNLL-U format (such as the training section of the English Web Treebank, provided with the lab),
- for each tree in the treebank, computes a projective dependency tree that minimises the total amount of lifting (as explained above), and
- outputs the transformed tree.

To achieve this goal you will have to implement the Eisner algorithm, including a backpointer mechanism for constructing the dependency tree with minimal lifting.

Your script should work in the same way as the script `projectivize.py` that comes with this lab. That script implements another algorithm for projectivisation; however, both algorithms produce the same results. To test your implementation, you can therefore compare the output of your script to the output of `projectivize.py` using a tool such as `diff`; the two outputs should be identical.